From Hard Direction Decoding to Time-Dependent Low-Dimensional Neural Decoding for Hand Trajectory Estimation

Yunxiang Cai, Sheng Jiang, Minghan Li, Haoqi Zhang
Department of Bioengineering, Imperial College London
Email: {y.cai25, sheng.jiang25, minghan.li25, haoqi.zhang25}@imperial.ac.uk

***Abstract*—This report evaluates two causal decoding frameworks on the BMI hand trajectory estimation task. The first employed an explicit dynamics-based framework combining SVM-based direction classification, velocity regression, and Kalman Filter to model movement evolution. The second adopted a low-dimensional decoding approach, using PCA, LDA, and kNN-based direction inference together with trajectory regression in the reduced feature space. To improve online stability, lightweight causal smoothing and trajectory blending were further incorporated. Results showed that the low-dimensional framework achieved more robust and accurate trajectory prediction than the more heavily parameterised dynamics-based approach. This suggests that, for this dataset, decoder performance depended more on matching the structure of the neural data than on increasing model complexity.**

***Index Terms*—Brain–Machine Interface, Neural Decoding, Low-Dimensional Representation, Time-Dependent Models, Causal Decoding, Trajectory Estimation**

I. Introduction

Brain–machine interfaces (BMIs) seek to transform neural activity into useful motor outputs, making continuous trajectory decoding a central problem in prosthetic control and motor neuroscience. In this task, the objective is to estimate the planar hand trajectory of a macaque from spike trains recorded from 98 neural units during reaches to eight target directions. The decoder must remain causal, beginning prediction at 320 ms and updating every 20 ms using only current and past observations.

A central challenge is choosing an appropriate level of model structure. Machine learning methods for neural decoding range from simple linear models to highly expressive nonlinear approaches, and method choice remains strongly task-dependent [1]. At the same time, motor cortical population activity often evolves on low-dimensional manifolds that capture substantial task-relevant structure [2]. This suggests that improved decoding may depend not only on model complexity, but also on how well the decoder matches the latent structure of the neural data.

Motivated by this trade-off, we first explored a *Hard Direction Decoder*, which performs early direction selection followed by dynamics-based trajectory estimation. We then developed a *Time-Dependent PCA Decoder*, which leverages low-dimensional representations throughout the decoding process and incorporates lightweight causal stabilisation. Our results suggest that, for this dataset, robust performance is achieved more effectively through appropriate low-dimensional structure and pipeline design than through increasing model complexity alone.

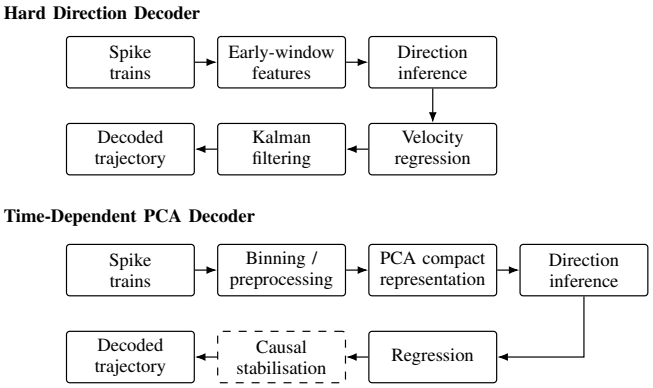

Fig. 1. Overview of the Hard Direction Decoder and the Time-Dependent PCA Decoder.

II. Methods

This section describes the two decoding frameworks explored in this study, summarised in Fig. 1. We first implemented a *Hard Direction Decoder*, motivated by the widespread use of staged direction inference and state-space models in neural trajectory estimation. We then developed a *Time-Dependent PCA Decoder*, which relies on compact low-dimensional representations and lightweight causal stabilisation for improved robustness.

A. Hard Direction Decoder

Our initial design followed a hard direction decoding paradigm, in which movement direction is inferred early and subsequent trajectory estimation is conditioned on a single selected direction. Because motor cortical neurons are known to exhibit directional tuning during reaching [3], the decoder first inferred movement direction from early neural activity using a supervised classifier, and then applied direction-specific continuous estimation. Among the direction-inference schemes explored during development, support vector machines (SVMs) were retained due to their empirical robustness.

Neural activity was converted into binned spike-count features, and recent spike history was used to predict two-dimensional velocity. To improve numerical stability in this regression stage, ridge regularisation was used:

$$\hat{W} = \arg\min_{W} \|XW - Y\|_2^2 + \lambda \|W\|_2^2, \tag{1}$$

where X denotes neural features, Y the target velocity, and λ the regularisation coefficient. The estimated velocity was then incorporated into a state-space model. Kalman filtering

was used to relate neural activity to movement kinematics [4], [5]. In its final form, the decoder employed a six-dimensional constant-acceleration state,

$$x_t = [p_x,\ p_y,\ v_x,\ v_y,\ a_x,\ a_y]^T, \tag{2}$$

with dynamics

$$x_t = Ax_{t-1} + w_t, \qquad z_t = Hx_t + v_t, \tag{3}$$

where w_t and v_t denote process and observation noise.

To better capture trajectory variability, noise statistics were learned from training data and adjusted for early and late decoding phases. Despite these refinements, the framework remained sensitive to errors introduced in the initial direction selection stage, motivating the development of a more robust decoding strategy.

B. Time-Dependent PCA Decoder

We next adopted a low-dimensional decoding strategy based on time-dependent neural representations. This design was motivated by the observation that motor cortical activity evolves on low-dimensional manifolds that capture task-relevant structure [2]. Rather than relying on a single global model, we constructed compact representations that adapt to the available neural evidence at each time point.

Spike trains were first converted into 20 ms bins, followed by a square-root transform and exponential moving average (EMA) smoothing to reduce noise. Features from all available causal time bins up to the current decoding time were concatenated to capture the temporal evolution of neural activity. Principal component analysis (PCA) was then applied to obtain a low-dimensional representation,

$$\tilde{x} = U_r^T(x - \mu), \tag{4}$$

where U_r contains the leading principal components.

The dimensionality r was not fixed manually, but determined automatically based on a variance retention criterion. Specifically, principal components were selected such that the cumulative explained variance exceeded a predefined threshold (varKeep).

Let $\{\lambda_i\}$ denote the eigenvalues sorted in descending order. The number of retained components r is chosen as the smallest integer satisfying

$$\frac{\sum_{i=1}^{r} \lambda_i}{\sum_{i=1}^{D} \lambda_i} \geq \text{varKeep}, \tag{5}$$

where D is the original feature dimension. In this work, the threshold was set to 0.85, allowing the representation to adapt to the intrinsic structure of the data while avoiding unnecessary high-dimensional noise.

In this reduced space, linear discriminant analysis (LDA) was used to enhance separation between movement directions, and direction inference was performed using distance-weighted k-nearest neighbours. This combination provided a compact, discriminative, and locally adaptive representation of neural activity.

Trajectory-related quantities were then estimated using direction-conditioned regression models in the reduced feature space. Compared with a single global mapping, this formulation enabled smoother and more stable neural-to-kinematic predictions while maintaining causality.

C. Lightweight Causal Stabilisation

Although the low-dimensional decoder was more robust, early-stage predictions could still exhibit instability due to directional uncertainty. We therefore introduced lightweight causal stabilisation mechanisms.

First, instead of relying on a single hard direction estimate, predictions from multiple candidate directions were softly combined:

$$\hat{y} = \sum_{m=1}^{M} w_m \hat{y}_m, \qquad \sum_{m=1}^{M} w_m = 1, \tag{6}$$

which reduced sensitivity to early misclassification. Second, predictions were blended with direction-specific mean trajectories, with the influence of the mean decaying over time.

All operations were strictly causal and relied only on past and current observations. These mechanisms did not replace the core decoder, but improved stability while preserving a compact and interpretable structure.

III. Results

We evaluated decoding performance using the root mean squared error (RMSE) of the reconstructed trajectories, together with runtime and the competition score. Table I summarises representative results from the Hard Direction Decoder, while Table II reports the performance of the proposed Time-Dependent PCA Decoder.

TABLE I
PERFORMANCE OF THE HARD DIRECTION DECODER (REPRESENTATIVE VARIANTS).

Method	RMSE	Time (s)	Score
Template + Ridge + Integration	21.82	1.57	19.79
Template + PCA + Ridge + Kalman	21.09	1.85	19.17
SVM + Ridge + Kalman	22.28	3.59	20.41
SVM + Ridge + learned CA-KF	22.09	2.61	20.14

Within the Hard Direction Decoder, improvements such as PCA-based denoising, alternative classifiers (e.g., SVM), and refined Kalman filtering resulted in only modest performance gains, with RMSE remaining around 21–23.

TABLE II
PERFORMANCE OF THE TIME-DEPENDENT PCA DECODER.

Method	RMSE	Time (s)	Score
Baseline time-dependent model	8.03	3.44	7.57
+ soft direction blending	7.83	3.42	7.38
Final model (tuned PCA + smoothing)	**7.50**	3.17	**7.07**

In contrast, the proposed Time-Dependent PCA Decoder achieved a substantially lower RMSE of approximately 7.5, representing a significant improvement over the Hard Direction Decoder. This indicates that redesigning the decoding

pipeline is more effective than refining individual components within a fixed structure.

The performance gain can be attributed to three main factors. First, PCA-based low-dimensional representations provide a compact and noise-robust description of neural activity. Second, direction inference in this reduced space improves reliability compared to early hard decisions. Third, soft blending of multiple candidate directions and lightweight causal stabilisation reduce sensitivity to early decoding errors and improve temporal consistency.

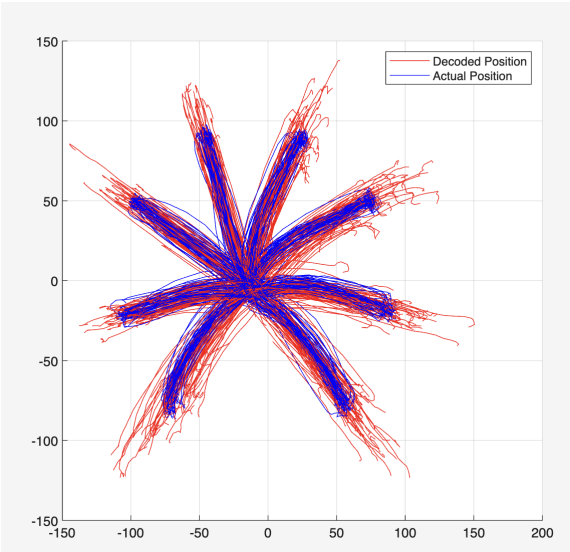

Fig. 2. Decoded trajectories (red) versus ground truth (blue) using the Hard Direction Decoder. The decoded paths capture the general reaching directions, but show larger deviations and greater trajectory spread, reflecting sensitivity to early hard direction selection and accumulated estimation error.

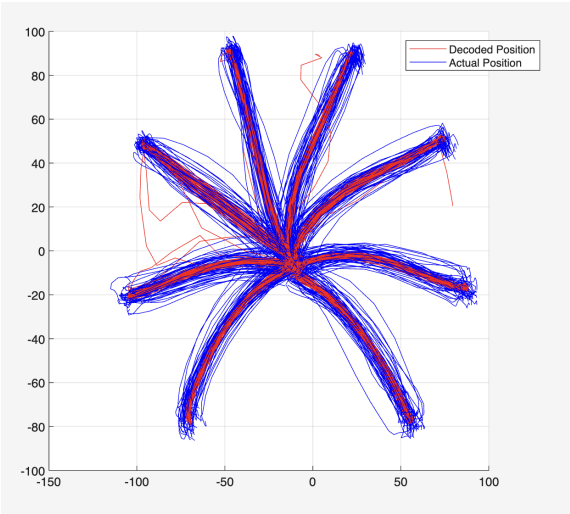

Fig. 3. Decoded trajectories (red) versus ground truth (blue) using the Time-Dependent PCA Decoder. Compared with the Hard Direction Decoder, the decoded trajectories are more tightly aligned with the true movements and show smoother behaviour across reaching directions.

Figure 2 shows that the Hard Direction Decoder captures the overall radial structure of the reaching task, but produces visibly larger deviations from the true trajectories. In contrast, Fig. 3 shows that the Time-Dependent PCA Decoder generates trajectories that are more tightly aligned with the ground truth, consistent with its lower RMSE.

Figure 3 illustrates representative decoded trajectories. The proposed method closely follows the ground truth across all reaching directions, while maintaining smooth and consistent trajectories over time.

Overall, these results demonstrate that appropriate use of low-dimensional structure and pipeline design leads to substantial improvements in decoding accuracy, while remaining fully causal and computationally efficient.

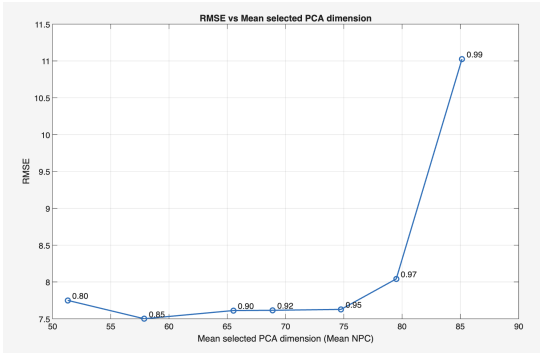

Fig. 4. RMSE as a function of the mean number of selected PCA dimensions. Each point corresponds to a different variance retention threshold. An intermediate dimensionality achieves the best performance, while too few or too many components lead to increased error.

We further analysed the effect of PCA dimensionality on decoding performance (Fig. 4). The results show a non-monotonic relationship between the number of retained principal components and RMSE.

Best performance is achieved at an intermediate dimensionality, corresponding to a variance retention of approximately 0.85. When too few components are retained, important task-relevant information is lost. In contrast, retaining too many components introduces noise and degrades performance.

This result highlights the importance of balancing information preservation and noise suppression when constructing low-dimensional neural representations.

IV. Discussion

The results of this study highlight that decoding performance is influenced more by the overall pipeline structure than by the choice of individual components. While the Hard Direction Decoder incorporated increasingly sophisticated elements, including SVM-based classification and Kalman filtering, these refinements resulted in only marginal improvements. This suggests that improving individual modules within a fixed architecture is insufficient when the underlying information flow is suboptimal.

A key limitation of the Hard Direction Decoder lies in its reliance on early hard direction selection. Errors made

at this stage propagate through subsequent velocity estimation and state-space filtering, leading to persistent trajectory deviations. In contrast, the proposed Time-Dependent PCA Decoder mitigates this issue by delaying commitment to a single direction and allowing multiple candidate directions to contribute through soft combination. This reduces sensitivity to early uncertainty and improves robustness in the causal decoding setting.

Another important factor is the use of time-dependent models. By constructing separate representations for different time points, the decoder adapts to the amount of available neural information as the movement unfolds. Early predictions rely on compact, noise-robust features, while later predictions benefit from richer temporal context. This contrasts with a static model, which must balance these competing requirements globally.

The use of PCA further supports this design by providing a low-dimensional representation that captures task-relevant neural structure while suppressing noise. Importantly, the dimensionality is selected automatically based on a variance retention criterion, avoiding manual tuning and enabling the representation to adapt to the intrinsic structure of the data. As shown in the results, decoding performance depends critically on this balance: overly aggressive dimensionality reduction leads to information loss, whereas retaining too many components introduces noise.

Finally, the proposed approach achieves these improvements without relying on complex state-space models. Instead, lightweight causal smoothing and feature-based temporal integration provide sufficient stability for accurate trajectory reconstruction. This demonstrates that, for this task, appropriate use of low-dimensional structure and pipeline design can outperform more complex dynamic models while remaining computationally efficient and fully causal.

Overall, these findings suggest that effective neural decoding depends not only on model complexity, but more importantly on how information is structured, represented, and propagated through the decoding pipeline.

V. Conclusion

In this work, we investigated neural trajectory decoding under a strictly causal setting and compared two decoding strategies: a Hard Direction Decoder and a Time-Dependent PCA Decoder. While the former relies on early direction classification followed by model-based trajectory estimation, the latter adopts a low-dimensional, time-dependent representation combined with lightweight stabilisation.

Our results show that substantial performance improvements are achieved not by increasing model complexity, but by redesigning the decoding pipeline. In particular, the use of time-dependent representations, PCA-based dimensionality reduction, and soft direction combination significantly improves robustness and reduces error propagation, leading to a marked reduction in RMSE.

These findings suggest that effective neural decoding depends critically on how information is represented and integrated over time. Future work may explore extending this framework to more complex movement tasks or incorporating adaptive representations that further capture nonlinear neural dynamics.

The implementation of the proposed methods is available at: https://github.com/shengj1ang/DanceMonkey.

Acknowledgment

We would like to express their sincere gratitude to Dr. Clopath Claudia, and all the GTAs for their valuable guidance, feedback, and support throughout this project.

Attributions

All team members actively and collaboratively contributed to the project.

Yunxiang Cai led the optimisation and final development of the Time-Dependent PCA Decoder.

Sheng Jiang contributed to the initial development of the Time-Dependent PCA Decoder and was responsible for data visualisation and figure generation.

Minghan Li developed the Hard Direction Decoder and explored its variants.

Haoqi Zhang contributed to the development of the Hard Direction Decoder and participated in the initial structuring and drafting of the report.

All members were involved in discussions, experimental design, result analysis, and refinement of the final report.

Artificial intelligence tools were used in a limited and supportive manner during the project. Specifically, they assisted with language polishing, grammar correction, and minor debugging suggestions. All core algorithm design, implementation decisions, and experimental evaluations were carried out by the team.

References

[1] J. I. Glaser, A. S. Benjamin, R. H. Chowdhury, M. G. Perich, L. E. Miller, and K. P. Kording, "Machine learning for neural decoding," *eNeuro*, vol. 7, no. 4, 2020.

[2] J. A. Gallego, M. G. Perich, L. E. Miller, and S. A. Solla, "Neural manifolds for the control of movement," *Neuron*, vol. 94, no. 5, pp. 978–984, 2017.

[3] A. P. Georgopoulos, J. F. Kalaska, R. Caminiti, and J. T. Massey, "On the relations between the direction of two-dimensional arm movements and cell discharge in primate motor cortex," *The Journal of Neuroscience*, vol. 2, no. 11, pp. 1527–1537, 1982.

[4] W. Wu, A. Shaikhouni, J. P. Donoghue, and M. J. Black, "Closed-loop neural control of cursor motion using a kalman filter," in *Proceedings of the 26th Annual International Conference of the IEEE Engineering in Medicine and Biology Society*, 2004, pp. 4126–4129.

[5] W. Q. Malik, W. Truccolo, E. N. Brown, and L. R. Hochberg, "Efficient decoding with steady-state kalman filter in neural interface systems," *IEEE Transactions on Neural Systems and Rehabilitation Engineering*, vol. 19, no. 1, pp. 25–34, 2011.